

Mobile ML Workshop

ML@B · Spring 2018



Installation and Setup

- `pip install numpy keras tensorflow`
- Download Xcode from the Mac App Store
- `git clone`
`https://github.com/mlberkeley/mobile_ml_workshop`
- **Sign In Here:** www.tinyurl.com/mobile-ml

Who We Are



Naren Krishna
ML@B VP of Resources
EECS
narenk [at] berkeley.edu



Gokul Swamy
ML@B Project Manager
EECS
gokul.swamy [at] berkeley.edu


What we will be building



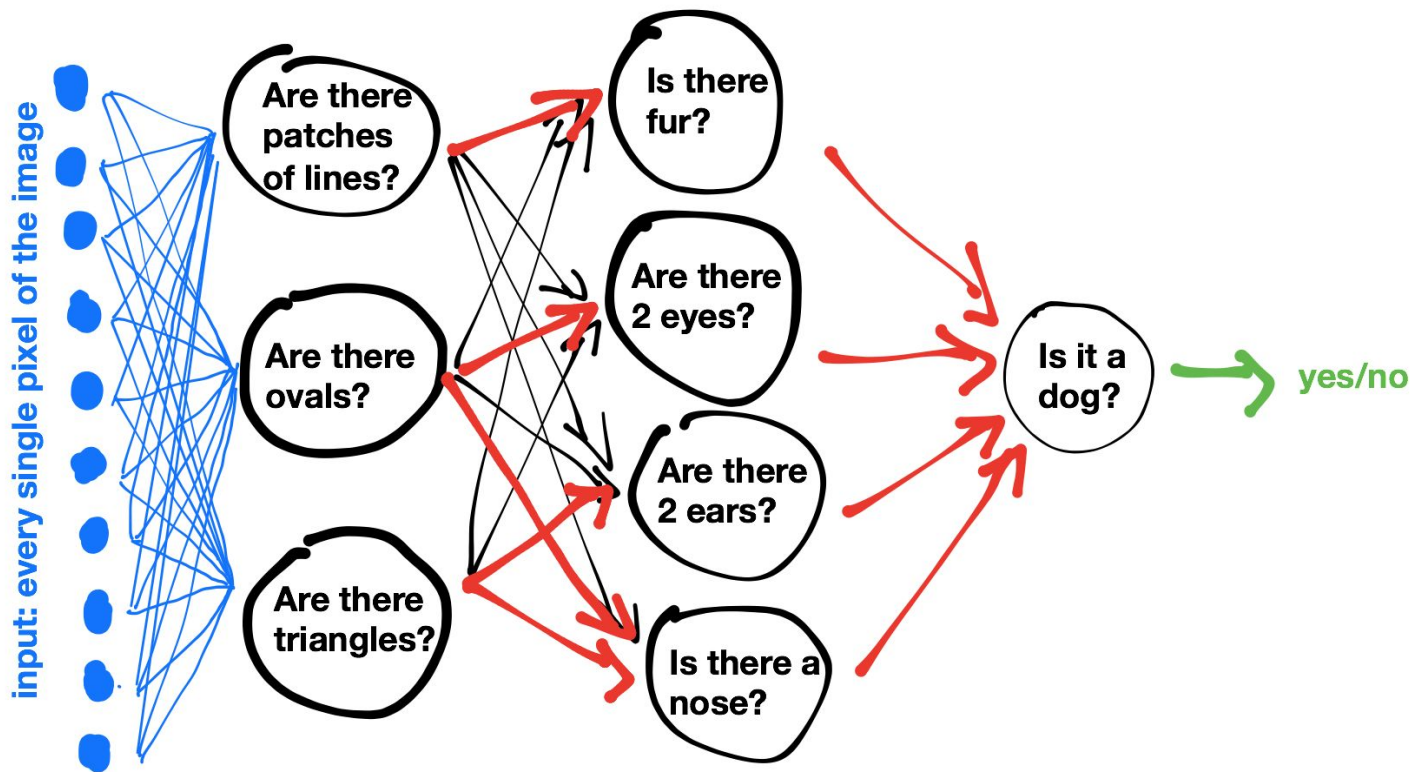
Agenda

- Deep Learning Overview
- Principles of Mobile ML
- Swift Overview
- CoreML
- Building a Mobile ML System


Quick Deep Learning Overview



Why are Neural Nets so Powerful?




Structure of Basic Feedforward Network




$$\hat{y} = f(\vec{x}) = \sigma \left(W^{(2)} \sigma \left(W^{(1)} \vec{x} + \beta^{(1)} \right) + \beta^{(2)} \right)$$

Gradient Descent

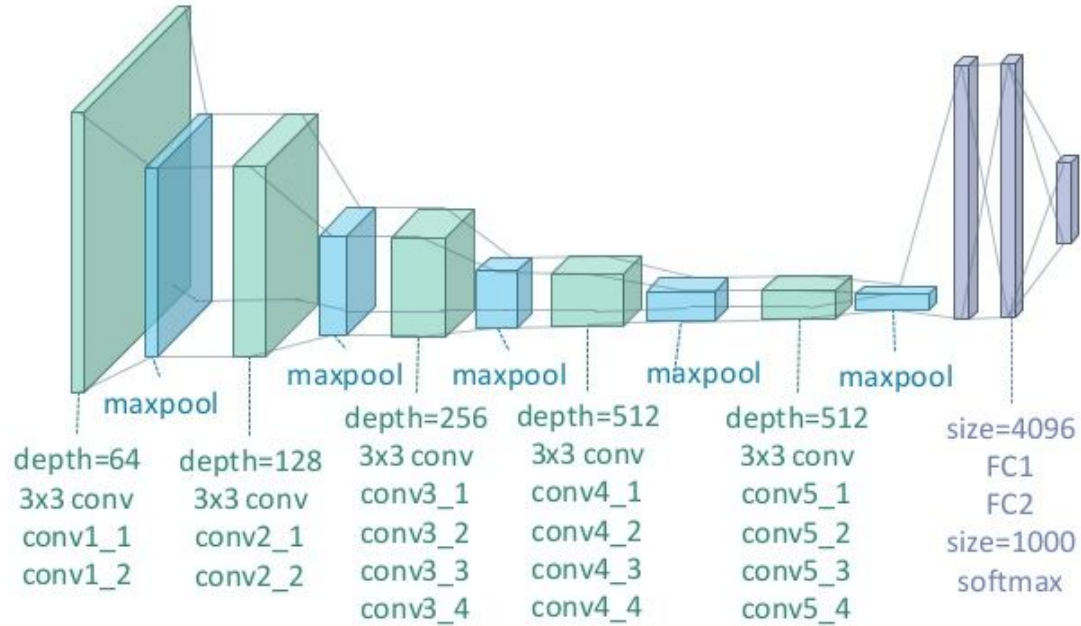
Define the cost function $J(w)$ to be the average of the $L(f(x), y)$ over all the training examples x . Note that the cost function $J(w)$ is a function of the weights of the network.



$$\vec{w} \leftarrow \vec{w} - \eta \frac{\partial J}{\partial \vec{w}}$$



Convolutional Neural Networks



Convolutions

Weight Filter

1	0	1
0	1	0
1	0	1

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature


Convolutions




*

1	2	1
0	0	0
-1	-2	-1


—




Convolutional Neural Networks



Early Layers Learn Edge Detectors



Later Layers Build Interesting Abstractions




Layer 2 of AlexNet



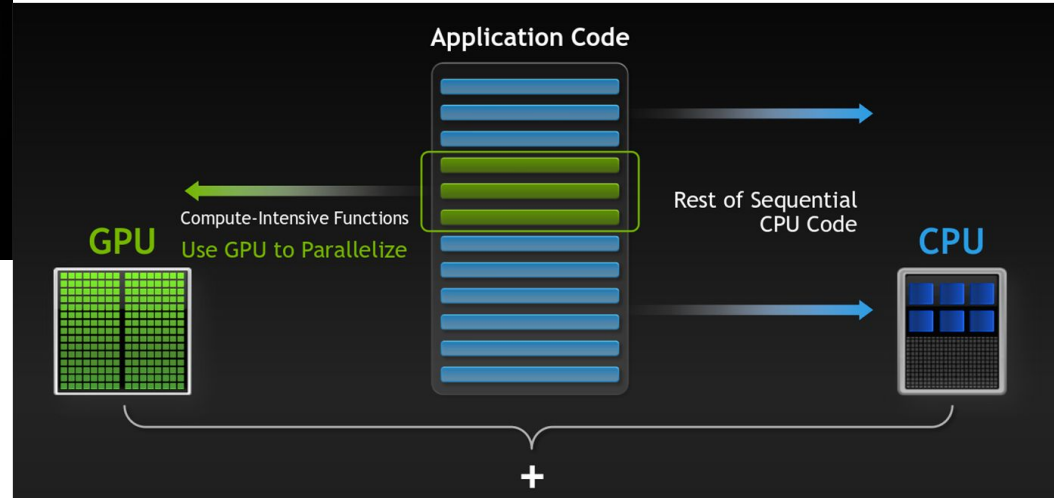
Layer 4 of AlexNet

Deep Learning for MNIST Digit Classification



[illegible]


How does Deep Learning leverage computational resources?



Why Mobile?



FUNDING BY ARTIFICIAL INTELLIGENCE CATEGORY




Challenges for Mobile ML

- Computational resource-limited environment
- Restrictions on amount of power we can consume
- Cannot assume connectivity to the internet or power source
- Very sensitive and personal data on people's phones.
- Data is not i.i.d (different amounts and sources of data)




Strategies to Deal With Challenges

- Balance between training on device (better for privacy) and training on a server (better for aggregation and compute resources)
- If training on device, be sure to only do it when user has enough battery and computational resources (for example, plugged in overnight)
- If training on server, be sure to anonymize data in some way (look into differential privacy if curious) and not wait on disconnected users to deliver data
- Federated Learning: Training on mobile and using a weighted average to update model on server



Swift Overview



What is Swift?

- Swift is the language currently used to develop iOS Apps
- ~10th most popular programming language
- Created by Apple, [open-sourced](#)
- Built on top of speedy LLVM (low-level virtual machine) compiler
- Supports Python-like REPL and runs on Linux

Constants/Variables

- Use **let** to declare a constant, **var** to declare a variable
- `//` to comment a line
- Swift will infer a type but you can explicitly define a type with a colon following the variable name

```
import Foundation
```

```
var x = "Hello"  
x = "World"
```

```
let implicitInt = 50  
let explicitInt: Int = 70
```

```
// error  
implicitInt = 94
```

Arrays

- Specified using square brackets around type
- Do not have a fixed size
- Use **.append** or **+=** to add elements
Use **.count** to get size of array
Swift supports enhanced for loops

```
var letters: [String] =  
    ["a", "b", "c", "d"]  
  
for letter in letters {  
    print(letter)  
}
```

Control Flow

- Use ... to specify an inclusive interval
- Use ..< to specify an upper-bound exclusive interval
- Also, "...\\(**var_name**)..." inserts variable into string

```
for i in 1..<4 {  
    print("I'm #\\(i)")  
}  
  
// Prints  
// I'm #1  
// I'm #2  
// I'm #3
```


Functions

- Explicit parameter and return types
- Interior and exterior parameter names

```
func remind(_ name: String, to action: String) -> String {  
    return "Remind \$(name) to \$(action)"  
}
```

```
// Called as:  
remind("Gokul", to: "sleep")
```

Closures

- Unnamed functions that can be passed as arguments to other functions
- Used for functional programming

```
// All Equivalent
func isAGoodBoy(animal: String) -> Bool {
    return animal == "dog"
}

let isAGoodBoy= {animal in
    return animal == "dog"
}

let isAGoodBoy= {
    return $0 == "dog"
}

let isAGoodBoy= {
    $0 == "dog"
}

// Functional Swift
let animals = ["dog", "cat", "fish", "fox"]
let goodBoys = animals.filter({$0 == "dog"})
```

Classes

- Use **self** to refer to current object
- Use **super** to reference superclass
- Use a colon to specify subclassing

```
class Animal {  
    private String name  
  
    init(name: String) {  
        self.name = name  
    }  
}  
  
class Dog: Animal {  
    init() {  
        super.init(name: "Fido")  
    }  
}
```

Protocols

- Similar to interfaces in other languages
- Can specify variables and functions
- Can specify whether variable should be gettable or settable
- Same colon syntax as subclassing

```
protocol DogOwner {  
    var dog: Dog {get set}  
    func pet() -> String  
}
```

```
class Sandy: DogOwner {  
    var dog: Dog = Dog()  
    func pet() -> String {  
        return "Woof!"  
    }  
}
```

Protocols

- Similar to interfaces in other languages
- Can specify variables and functions
- Can specify whether variable should be gettable or settable
- Same colon syntax as subclassing

```
protocol DogOwner {  
    var dog: Dog {get set}  
    func pet() -> String  
}
```

```
class Sandy: DogOwner {  
    var dog: Dog = Dog()  
    func pet() -> String {  
        return "Woof!"  
    }  
}
```

Optionals

- One of the signature features (and sometimes headaches) of Swift
- A type that indicates that an object of another (non-optional) type exists or a nil value
- Unwrap to get at data that is potentially contained

Optionals: Syntax

```
var greeting: String? = "Hello World" // Use ? to declare optional
```


```
print(greeting!) // Use ! to forcibly unwrap
```

```
greeting = nil  
print(greeting!) // Causes Error
```

```
if let unwrapped = greeting { // Safe Unwrapping  
    print(unwrapped) // Only reached if non-nil value  
}
```

```
// Optional chaining  
print(greeting?.count) // (checks and sees if greeting is  
                        // nil, avoiding error)
```

```
// Use for early exit and to avoid pyramid of doom  
guard let unwrapped = greeting else {return}
```



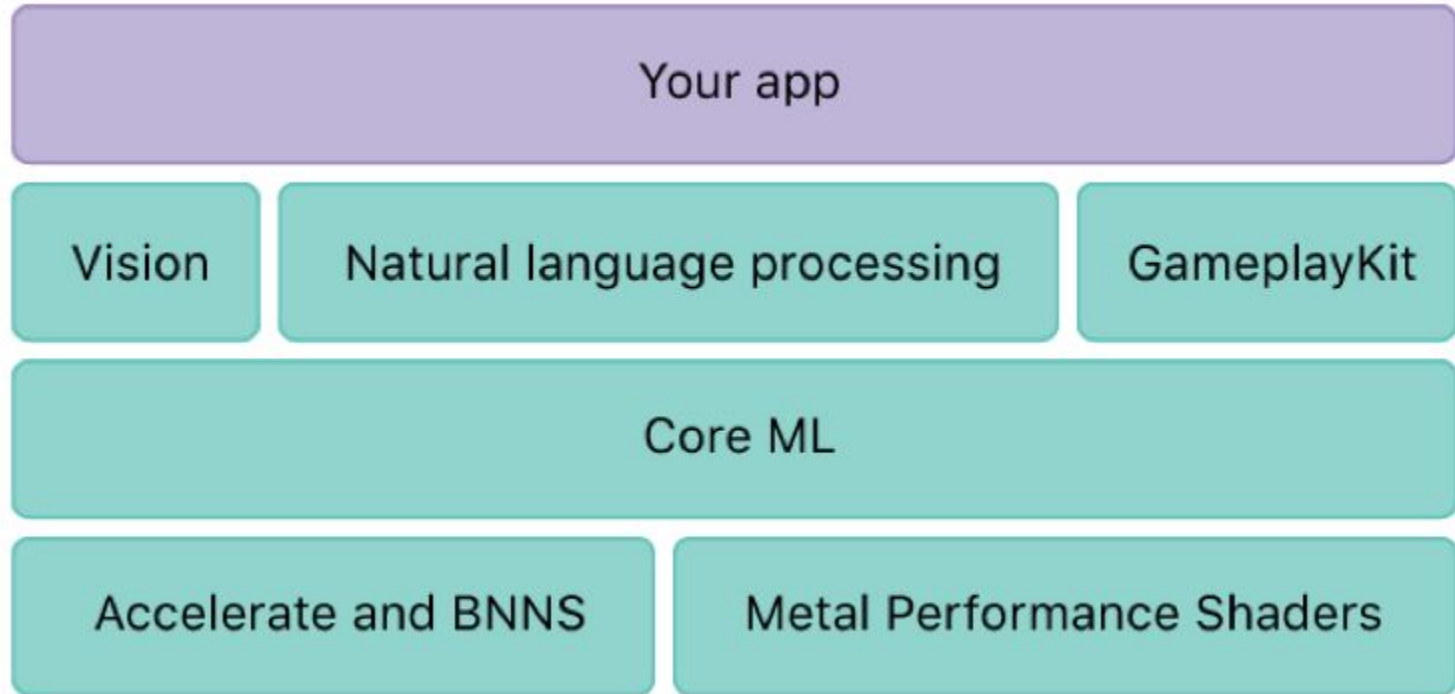
What is CoreML?

- CoreML released at WWDC 2017
- Represents a shift in traditional developer model
 - Researcher creates model
 - Apple takes care of hard part of implementation
 - You just have to connect the above

CoreML

- Lets you convert ML code written in Python into a **.mlmodel** file can be dropped directly into your app
- Hard part: optimizes code for iOS devices
- Some helpful preprocessing included
- Does not currently support on-device training

Architecture



Vision

- Allows you to detect face landmarks, perform OCR, scan barcodes, perform optical flow
- Slower but more accurate than CoreImage and AVFoundation
- Three Steps
 1. Call **VNRequestHandler**
 2. Which executes a **VNRequest**
 3. And returns some **VNObservation** for you to process

Vision: Code

- Be sure to add Camera Usage Description to Info.plist and import Vision

```
func detectText(image: UIImage){
    let textRequest = VNDetectTextRectanglesRequest(completionHandler: self.detectTextCompletionHandler)
    let textRequestHandler = VNImageRequestHandler(cgImage: image.cgImage!, options: [:])
    do {
        try textRequestHandler.perform([textRequest])
    } catch {
        print(error)
    }
}

func detectTextCompletionHandler(request: VNRequest, error: Error?){
    guard let results = request.results as? [VNTextObservation] else {return}
    var boxes = [VNRectangleObservation]()
    for result in results {
        if let characterBoxes = result.characterBoxes {
            for box in characterBoxes {
                boxes.append(box)
            }
        }
    }
    // do something here with boxes
}
```

NSLinguisticTagger

- Part of NS but revamped with deep learning for 2017
- On device processing so no privacy worries
- Can identify language, tokenize (split up into chunks), lemmatize (give root form of word), and detect named entities (rigid designators of concepts that may be important)
- Two Steps
 - Create **NSLinguisticTagger** with tagSchemes set to application
 - Set **tagger.string** to the text to be analyzed.

NSLinguisticTagger: Code

- For example, to detect dominant language:

```
let tagger = NSLinguisticTagger(tagSchemes: [.language], options: 0)
tagger.string = "NSLinguisticTagger provides text processing APIs."
if let language = tagger.dominantLanguage {
    print(language)
}
```

- For all use case sample code, check out [this link](#)



Building a Mobile ML System




Setup

- We'll be using Python 3.5
- `git clone`
`https://github.com/mlberkeley/mobile_ml_workshop`
- `pip install numpy keras tensorflow`
- Download Xcode from the Mac App Store
- `pip install h5py coremltools`
- `xcode-select --install`

Federated Learning

- Compute gradients based on individual user data
- Send gradients to server and update model by averaging gradients across users
- Maintains user data privacy



Algorithm 1 FederatedAveraging. The K clients are indexed by k ; B is the local minibatch size, E is the number of local epochs, and η is the learning rate.

Server executes:

initialize w_0

for each round $t = 1, 2, \dots$ **do**

$m \leftarrow \max(C \cdot K, 1)$

$S_t \leftarrow$ (random set of m clients)

for each client $k \in S_t$ **in parallel do**

$w_{t+1}^k \leftarrow \text{ClientUpdate}(k, w_t)$

$w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$

ClientUpdate(k, w): *// Run on client k*

$\mathcal{B} \leftarrow$ (split \mathcal{P}_k into batches of size B)

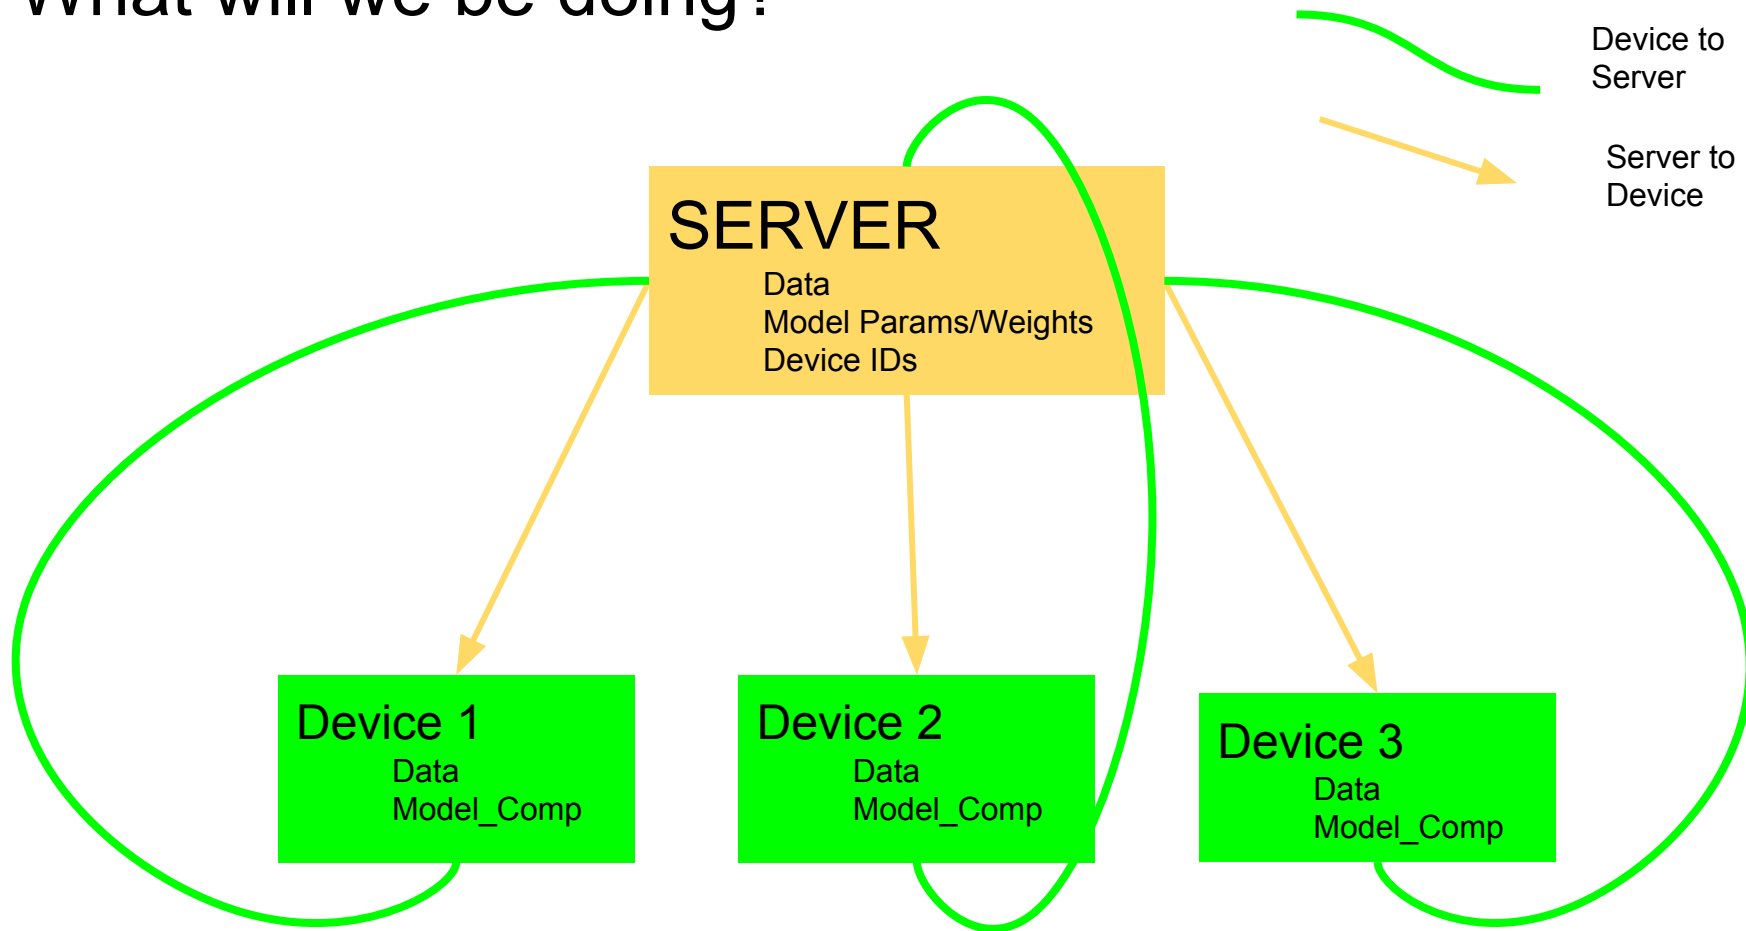
for each local epoch i from 1 to E **do**

for batch $b \in \mathcal{B}$ **do**

$w \leftarrow w - \eta \nabla \ell(w; b)$



return w to server

What will we be doing?



How can we make our model more accurate?

- Do preprocessing on MNIST images to make sure no pixels have gray values
 - Similar format as inputs from phone
- Change pen style on phone to sense pressure when writing
 - Makes digits written on phone more similar to MNIST images



Extra Resources



How to Deep Learn at Home

Run on CPU (really slow)

Buy a GPU (more expensive option)

Use AWS! - [Here's a Guide](#)

(Jan. 2018) OR

Use [Google Colab!](#) ([GPU guide](#))

Deep Learning Resources

[ML@B's Machine Learning Crash Course: Part 3](#)

[Stanford's CS231n Convolutional Networks for Visual Recognition](#)

[Deep Learning Textbook by Goodfellow et. al.](#)

[ML@B's Deep Learning Decal](#)

[Neural Networks and Deep Learning \(Backprop section\)](#)